AMAN VERMA

AI Engineer — Machine Learning / Deep Learning / NLP

aman.verma@example.com | +92 300 1234567 | Karachi, Pakistan | github.com/aman-verma | linkedin.com/in/aman-verma

SAMPLE — FICTIONAL / FOR PRACTICE ONLY. DO NOT USE TO MISREPRESENT.

Summary

Highly motivated AI Engineer with 4+ years of experience designing, developing, and deploying machine learning and deep learning solutions. Proven track record in NLP, computer vision, and scalable model productionization. Strong software engineering foundations and cloud deployment experience.

Key Skills

Languages: Python, SQL, Bash

Frameworks: PyTorch, TensorFlow, scikit-learn

Tools: Docker, Kubernetes, Airflow, MLflow, Git

Cloud: AWS (S3, EC2, SageMaker), GCP (Cloud Storage, AI Platform)

Other: Data engineering, model optimization, A/B testing, MLOps

Professional Experience

Senior AI Engineer — Nimbus Labs — Karachi, PK *(Jan 2023 — Present)*
• Led design and deployment of a real-time NLP pipeline for customer support; reduced average response triage time by 45%.
• Built transformer-based intent classification model (RoBERTa variant) achieving 93% F1 on production data.
• Designed CI/CD model deployment workflows using Docker, Kubernetes and MLflow, enabling safe model rollbacks and monitoring.

AI Engineer — DataForge Solutions — Remote *(Aug 2020 — Dec 2022)*
• Developed computer vision models for defect detection using EfficientNet; improved detection accuracy by 27% over baseline.
• Implemented data labeling tooling and active learning loop that reduced annotation costs by 60%.
• Collaborated with product and ML Ops teams to scale training to distributed GPU clusters.

Machine Learning Engineer (Intern → Jr.) — OpenMetrics Research *(Jun 2019 — Jul 2020)*
• Built end-to-end ML workflows for time-series forecasting using Prophet and LSTM models.
• Produced model evaluation dashboards and automated model retraining pipelines.

Projects

SmartAssist — Contextual Conversational Agent — Designed an end-to-end conversational agent with retrieval-augmented generation (RAG). Implemented semantic search over vector store and fine-tuned T5 for domain responses. Deployed as microservice with REST API.

AutoDefect — Visual Defect Detection — Trained EfficientDet-based pipeline for manufacturing defect detection, implemented test-time augmentation and model ensembling; achieved 0.88 mAP in production.

ML Monitoring Framework — Built monitoring for dataset drift and model performance using streaming metrics and alerting; integrated with Grafana and Prometheus.

Education

M.S. in Computer Science (Machine Learning), National Institute of Technology — 2019

B.S. in Computer Science, City University — 2017

Certifications & Publications

• AWS Certified Machine Learning – Specialty (2022)
• DeepLearning.AI TensorFlow Developer (2021)
• Verma, A., et al. (2022). 'Efficient Transformer Pruning for Low-Latency Inference.' Workshop on Practical ML.

Tools & Technologies

Python, PyTorch, TensorFlow, scikit-learn, Docker, Kubernetes, MLflow, AWS, GCP, SQL, Airflow

Interests

Open-source contribution, mentoring junior engineers, competitive programming, hiking.